

Introduction to Algorithms

Lecture 6 Greedy Method

Xue Chen

xuechen1989@ustc.edu.cn

2025 spring in



Outline

- 1 Introduction
- 2 Revisit Interval Scheduling
- 3 Max Interval Sum
- 4 Huffman Codes
- 5 Set Cover
- 6 Scheduling

Overview

Greedy method — a powerful algorithm design technique

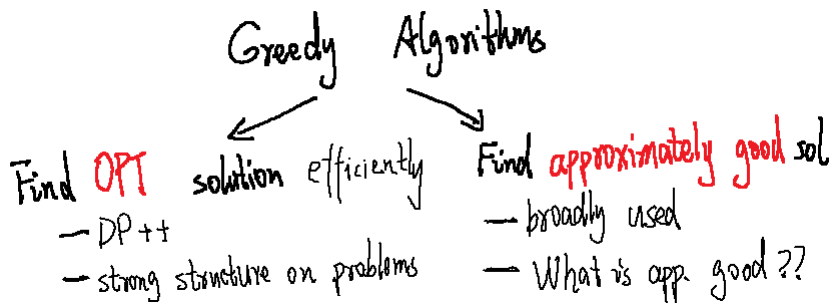
- 1 Very efficient
- 2 Simple to implement
- 3 Widely used in practice

Overview

Greedy method — a powerful algorithm design technique

- 1 Very efficient
- 2 Simple to implement
- 3 Widely used in practice

Two ways to understand it



Short Intro

For optimization problems, a greedy algorithm makes the choice that is the best at this moment — called local optimal

Short Intro

For optimization problems, a greedy algorithm makes the choice that is the best at this moment — called local optimal

- 1 For many problems, locally optimal choices lead to a globally optimal solution
- 2 For more problems, it does not — but one could show the solution is not bad

Short Intro

For optimization problems, a greedy algorithm makes the choice that is the best at this moment — called local optimal

- 1 For many problems, locally optimal choices lead to a globally optimal solution
- 2 For more problems, it does not — but one could show the solution is not bad

Main Focus

Proving the correctness of greedy algorithms is highly non-trivial.

Outline


- 1 Introduction
- 2 Revisit Interval Scheduling
- 3 Max Interval Sum
- 4 Huffman Codes
- 5 Set Cover
- 6 Scheduling

Interval Scheduling Problem

Recall: DP solves **weighted** scheduling


Description

- 1 n jobs in $[0, T]$
- 2 Job j starts from $s(j)$ and finishes at $f(j)$ with **weight** w_j
- 3 Two jobs are **compatible** if they don't overlap.
- 4 Goal: find a max-weight subset with compatible jobs.




Special Case: Unit Weights

What if all weights are 1?



Special Case: Unit Weights

What if all weights are 1?




Questions

While DP still works, faster or simpler algorithms?


Intuition

Let us try greedy choices.



Intuition


Let us try greedy choices.



① What if we pick the job with the earliest starting time?

Intuition


Let us try greedy choices.



- 1 What if we pick the job with the earliest starting time?
- 2 What if we pick the job with the earliest finish time?

Idea

Intuition: When there are multiple jobs, picking the one with the earliest finish time leaves the most space for future jobs.



Formal Proof

Theorem 16.1 in CLRS

For any problem S of job scheduling, let a_1 be the job with the earliest finish time. Then a_1 is included in some max-size compatible subset OPT of S .

Formal Proof

Theorem 16.1 in CLRS

For any problem S of job scheduling, let a_1 be the job with the earliest finish time. Then a_1 is included in some max-size compatible subset OPT of S .

- 1 Moreover, $OPT \setminus \{a_1\}$ is optimal for $[0, T] \setminus [0, f(a_1)]$.
- 2 The solution of subproblem $OPT' \cup \{a_1\}$ is optimal for S — by dynamic programming
- 3 Guarantee that we will reach a global OPT solution by making local OPT choice.

Algorithm

- 1 Recall $s[j]$ and $f[j]$ are the start and finish time separately
- 2 Sort all jobs according to the finish time $f[j]$.

GREEDY-ACTIVITY-SELECTOR(s, f)

```
1   $n = s.length$ 
2   $A = \{a_1\}$ 
3   $k = 1$ 
4  for  $m = 2$  to  $n$ 
5      if  $s[m] \geq f[k]$ 
6           $A = A \cup \{a_m\}$ 
7           $k = m$ 
8  return  $A$ 
```

Algorithm

- 1 Recall $s[j]$ and $f[j]$ are the start and finish time separately
- 2 Sort all jobs according to the finish time $f[j]$.

GREEDY-ACTIVITY-SELECTOR(s, f)

```
1   $n = s.length$ 
2   $A = \{a_1\}$ 
3   $k = 1$ 
4  for  $m = 2$  to  $n$ 
5      if  $s[m] \geq f[k]$ 
6           $A = A \cup \{a_m\}$ 
7           $k = m$ 
8  return  $A$ 
```

Analysis


Correctness follows from Theorem 16.1

Running time: $O(n \log n)$.

Discussion

While the running time is the same, greedy algorithm is arguably simpler than DP.

- 1 Easy to implement
- 2 However, it takes more effort to prove its correctness



Outline

- 1 Introduction
- 2 Revisit Interval Scheduling
- 3 Max Interval Sum
- 4 Huffman Codes
- 5 Set Cover
- 6 Scheduling

Problem

Given a sequence a_1, \dots, a_n of integers, find the largest sum of a consecutive interval, i.e., $\max_{k < \ell} \left\{ \sum_{i=k}^{\ell} a_i \right\}$.

$n=10$: 1 -2 3 -1 4 -2 3 | -5 4

Problem

Given a sequence a_1, \dots, a_n of integers, find the largest sum of a consecutive interval, i.e., $\max_{k < \ell} \left\{ \sum_{i=k}^{\ell} a_i \right\}$.

$n=10$: 1 -2 3 -1 4 -2 3 | -5 4

First idea

Enumerate all pairs k and ℓ — $O(n^2)$

1st Algorithm

```
function MAX-SUM(a)  
  ans = 0  
  for  $k = 1, \dots, n$  do  
    sum = 0  
    for  $\ell = k, \dots, n$  do  
      sum = sum +  $a[\ell]$   
      ans =  $\max\{ans, sum\}$   
  
  Return ans
```

Question

Can we design faster algorithms?

2nd Idea:

One standard trick: let $s[0] = 0$ and $s[i] = a_1 + a_2 + \cdots + a_i$ s.t. $\sum_{i=k}^{\ell} a_i = s[\ell] - s[k-1]$.

Problem becomes $\max_{k \leq \ell} \{s[\ell] - s[k-1]\}$.

2nd Idea:

One standard trick: let $s[0] = 0$ and $s[i] = a_1 + a_2 + \dots + a_i$ s.t. $\sum_{i=k}^{\ell} a_i = s[\ell] - s[k-1]$.

Problem becomes $\max_{k \leq \ell} \{s[\ell] - s[k-1]\}$.

Dynamic Programming

Let $f[\ell]$ denote the max-sum whose endpoint is ℓ :

$$f[\ell] = s[\ell] - \min_{k \leq \ell} s[k-1]$$

$$ans = \max \{f[1], \dots, f[n]\}$$

2nd Idea:

One standard trick: let $s[0] = 0$ and $s[i] = a_1 + a_2 + \dots + a_i$ s.t. $\sum_{i=k}^{\ell} a_i = s[\ell] - s[k-1]$.

Problem becomes $\max_{k \leq \ell} \{s[\ell] - s[k-1]\}$.

Dynamic Programming

Let $f[\ell]$ denote the max-sum whose endpoint is ℓ :

$$f[\ell] = s[\ell] - \min_{k \leq \ell} s[k-1]$$

$$ans = \max \{f[1], \dots, f[n]\}$$

Time $O(n \log n)$: Use a data structure (e.g., heap or BST) to find $\min_{k \leq \ell} \{s[k-1]\}$ for every ℓ in time $O(\log n)$.

Improve DP via greedy method

Back to $f[\ell] = s[\ell] - \min_{k < \ell} s[k - 1]$, consider ℓ and $\ell + 1$:

$$f[\ell] = s[\ell] - \min_{k \leq \ell} s[k - 1],$$

$$f[\ell + 1] = s[\ell + 1] - \min_{k \leq \ell + 1} s[k - 1]$$

Improve DP via greedy method

Back to $f[\ell] = s[\ell] - \min_{k < \ell} s[k - 1]$, consider ℓ and $\ell + 1$:

$$f[\ell] = s[\ell] - \min_{k \leq \ell} s[k - 1],$$

$$f[\ell + 1] = s[\ell + 1] - \min_{k \leq \ell + 1} s[k - 1]$$

OBS on \min_k

$$\min_{k \leq \ell + 1} s[k - 1] = \min \left(\min_{k \leq \ell} s[k - 1], s[\ell] \right)$$

Improve DP via greedy method

Back to $f[\ell] = s[\ell] - \min_{k < \ell} s[k - 1]$, consider ℓ and $\ell + 1$:

$$f[\ell] = s[\ell] - \min_{k \leq \ell} s[k - 1],$$

$$f[\ell + 1] = s[\ell + 1] - \min_{k \leq \ell + 1} s[k - 1]$$

OBS on \min_k

$$\min_{k \leq \ell + 1} s[k - 1] = \min \left(\min_{k \leq \ell} s[k - 1], s[\ell] \right)$$

function MAX-SUM(a)

$ans = 0, k = 0$

Compute $s[1], \dots, s[n]$

for $\ell = 1, \dots, n$ **do**

$k = \arg \min \{s[k], s[\ell - 1]\}$

$ans = \max\{ans, s[\ell] - s[k]\}$

Return ans

Greedy Algorithm

Actually, it is a greedy algorithm while it looks like a DP.

An alternate implementation

Let sum_ℓ denote the largest sum whose endpoint is ℓ s.t.

$$sum_\ell = \max\{a_\ell, sum_{\ell-1} + a_\ell\}.$$

function MAX-SUM(a)

$ans = 0, sum = 0$

for $\ell = 1, \dots, n$ **do**

$sum = \max\{0, sum\} + a_\ell$

$ans = \max\{ans, sum\}$

Return ans

Summary

- ① Greedy algorithms are elegant and fast
- ② Need more analysis and proofs (compare to DP and divide&conquer)

Summary

- 1 Greedy algorithms are elegant and fast
- 2 Need more analysis and proofs (compare to DP and divide&conquer)
- 3 So far, view it as a way to improve DP
- 4 Next, greedy algorithms solve many problems where DP can't help with

Summary

- 1 Greedy algorithms are elegant and fast
- 2 Need more analysis and proofs (compare to DP and divide&conquer)
- 3 So far, view it as a way to improve DP
- 4 Next, greedy algorithms solve many problems where DP can't help with

Strongly recommend: Read 15.3 and 16.2 to understand DP and greedy algorithms better

Outline

- 1 Introduction
- 2 Revisit Interval Scheduling
- 3 Max Interval Sum
- 4 Huffman Codes
- 5 Set Cover
- 6 Scheduling

Introduction

A classical problem in data compression

— solved by Huffman at 1952 in MIT

- 1 Consider a data set with alphabet say $\{a, b, c, \dots, z\}$, where each character appears f_a, \dots, f_z times
- 2 Find a binary encoding of $\{a, b, c, \dots, z\}$, called codewords, to minimize the total length.

	$a \times 85$	$b \times 50$	$c \times 60$
encode	$a \ 0$ $b \ 00$ $c \ 1$	$a \ 0$ $b \ 10$ $c \ 11$	$a \ 00$ $b \ 01$ $c \ 1$
	invalid!	length $1 \times 85 + 2 \times 50 + 2 \times 60$	$2 \times 85 + 2 \times 50 + 1 \times 60$

Introduction

A classical problem in data compression

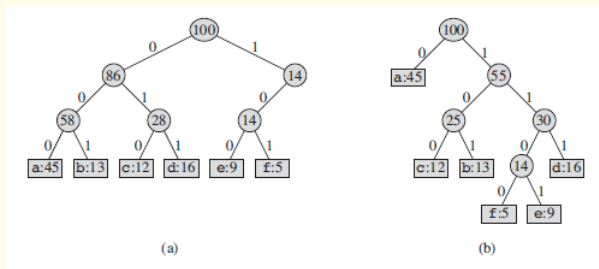
— solved by Huffman at 1952 in MIT

- 1 Consider a data set with alphabet say $\{a, b, c, \dots, z\}$, where each character appears f_a, \dots, f_z times
- 2 Find a binary encoding of $\{a, b, c, \dots, z\}$, called codewords, to minimize the total length.
- 3 Extra requirement — prefix codes: no codeword is a prefix of some other codewords.

	$a \times 85$	$b \times 50$	$c \times 60$
encode	$a \ 0$ $b \ 00$ $c \ 1$	$a \ 0$ $b \ 10$ $c \ 11$	$a \ 00$ $b \ 01$ $c \ 1$
	invalid!	length $1 \times 85 + 2 \times 50 + 2 \times 60$	$2 \times 85 + 2 \times 50 + 1 \times 60$

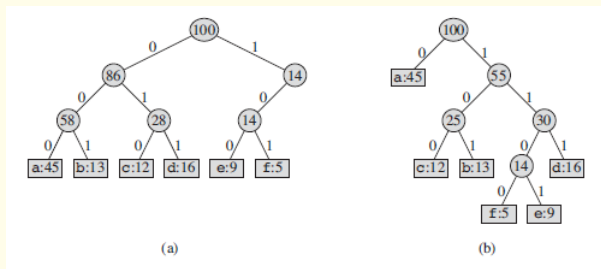
Idea: Encode them as a tree

- 1 Each leaf contains a character, whose codeword is the $\{0, 1\}$ -path from the root
- 2 The number on each node denotes the sum of frequencies of all leaves in its subtree



Idea: Encode them as a tree

- 1 Each leaf contains a character, whose codeword is the $\{0, 1\}$ -path from the root
- 2 The number on each node denotes the sum of frequencies of all leaves in its subtree



OBS

- 1 It satisfies the prefix codes requirement.
- 2 The total length is the sum of frequencies on nodes.

Greedy Algorithm

New goal: Construct a tree minimize the sum of weights

- 1 Let C be the set of n characters
- 2 For each $c \in C$, $c.freq$ is its frequency

(a)




f:5	e:9	c:12	b:13	d:16	a:45
-----	-----	------	------	------	------

Greedy Algorithm

New goal: Construct a tree minimize the sum of weights

- 1 Let C be the set of n characters
- 2 For each $c \in C$, $c.freq$ is its frequency
- 3 Reduce to subproblem: Merge two nodes with the lowest frequencies into one

(a) f:5 e:9 c:12 b:13 d:16 a:45







Greedy Algorithm

New goal: Construct a tree minimize the sum of weights

- 1 Let C be the set of n characters
- 2 For each $c \in C$, $c.freq$ is its frequency
- 3 Reduce to subproblem: Merge two nodes with the lowest frequencies into one

(a) f:5 e:9 c:12 b:13 d:16 a:45



Algorithm Description

Implement it via a min-heap

```
HUFFMAN( $C$ )
1   $n = |C|$ 
2   $Q = C$ 
3  for  $i = 1$  to  $n - 1$ 
4      allocate a new node  $z$ 
5       $z.left = x = \text{EXTRACT-MIN}(Q)$ 
6       $z.right = y = \text{EXTRACT-MIN}(Q)$ 
7       $z.freq = x.freq + y.freq$ 
8       $\text{INSERT}(Q, z)$ 
9  return  $\text{EXTRACT-MIN}(Q)$     // return the root of the tree
```

Algorithm Description

Implement it via a min-heap

```
HUFFMAN(C)
1   $n = |C|$ 
2   $Q = C$ 
3  for  $i = 1$  to  $n - 1$ 
4      allocate a new node  $z$ 
5       $z.left = x = \text{EXTRACT-MIN}(Q)$ 
6       $z.right = y = \text{EXTRACT-MIN}(Q)$ 
7       $z.freq = x.freq + y.freq$ 
8       $\text{INSERT}(Q, z)$ 
9  return  $\text{EXTRACT-MIN}(Q)$     // return the root of the tree
```

Analysis

- 1 Running Time: $O(n \log n)$ by heap
- 2 Correctness: Next 2 slides

Correctness

2 Steps: Consider one **merge** operation

- 1 \exists a optimal solution T with the **merged** pattern
- 2 For the **subproblem after merging**, its optimal solution $T' \Rightarrow$ an optimal solution T of original problem

Correctness

2 Steps: Consider one **merge** operation

- 1 \exists a optimal solution T with the **merged** pattern
- 2 For the **subproblem after merging**, its optimal solution $T' \Rightarrow$ an optimal solution T of original problem

Step 1: Lemma 16.2 in CLRS

Let $C = \{c_1, \dots, c_n\}$ where $c_1.freq \geq c_2.freq \geq \dots \geq c_n.freq$. Then \exists an **optimal** prefix code whose binary tree T has a node with children c_{n-1} and c_n as two of the deepest leaves.

Correctness cont.

Step 2: Lemma 16.3 in CLRS

Let $C' = \{c_0, c_1, \dots, c_{n-2}\}$ where $c_0.freq = c_{n-1}.freq + c_n.freq$. Let T' denote the binary tree of any optimal encoding of C' . Then the binary tree T of C , obtained by splitting the leaf c_0 into c_{n-1} and c_n , is optimal.

Correctness cont.

Step 2: Lemma 16.3 in CLRS

Let $C' = \{c_0, c_1, \dots, c_{n-2}\}$ where $c_0.freq = c_{n-1}.freq + c_n.freq$. Let T' denote the binary tree of any optimal encoding of C' . Then the binary tree T of C , obtained by splitting the leaf c_0 into c_{n-1} and c_n , is optimal.

- 1 Lemma 16.2 says it is OK to combine c_{n-1} and c_n since there exists an optimal solution doing that
- 2 Lemma 16.3 says any optimal solution of the subproblem is good

Summary


- 1 Greedy algorithms are elegant and fast
- 2 Proving Correctness is involved

Summary

- 1 Greedy algorithms are elegant and fast
- 2 Proving Correctness is involved
- 3 For many problems, greedy algorithms improves the running time of dynamic programming
- 4 Greedy algorithms provide efficient solutions for many problems where DP can not help with like Huffman code — more examples from computational hard problems

Summary

- 1 Greedy algorithms are elegant and fast
- 2 Proving Correctness is involved
- 3 For many problems, greedy algorithms improves the running time of dynamic programming
- 4 Greedy algorithms provide efficient solutions for many problems where DP can not help with like Huffman code — more examples from computational hard problems



Outline

- 1 Introduction
- 2 Revisit Interval Scheduling
- 3 Max Interval Sum
- 4 Huffman Codes
- 5 Set Cover
- 6 Scheduling

Approximation Algorithms

- ① Hard or slow to find optimal solutions in many problems — slow means a large polynomial time or super-polynomial time
- ② Computer Science is happy with an efficient algorithm if its output is **not bad**
- ③ Formally, it means approximately good — say the minimum is OPT , its answer is at most $\alpha \cdot OPT$ for some $\alpha > 1$
- ④ $\alpha :=$ approximation ratio
- ⑤ Finding approximation solutions is a central idea in CS — streaming algorithms (heavy hitters), big data algorithms, machine learning algorithms, privacy, cryptography, . . .

Set Cover

Given a ground set $[n] := \{1, 2, \dots, n\}$ and m subsets S_1, \dots, S_m , find the smallest selection of subsets S_i whose union is $[n]$.

Set Cover

Given a ground set $[n] := \{1, 2, \dots, n\}$ and m subsets S_1, \dots, S_m , find the smallest selection of subsets S_i whose union is $[n]$.

Greedy Algorithm

Repeat until $[n]$ is covered:

Pick S_i with the largest uncovered elements.

Set Cover

Given a ground set $[n] := \{1, 2, \dots, n\}$ and m subsets S_1, \dots, S_m , find the smallest selection of subsets S_i whose union is $[n]$.

Greedy Algorithm

Repeat until $[n]$ is covered:

Pick S_i with the largest uncovered elements.

- 1 Simple and easy to implement
- 2 Can not guarantee that the answer is optimal

Main Results

Theorem

If the optimal solution uses k subsets, the greedy algorithm uses at most $k \ln n$ subsets.

Main Results

Theorem

If the optimal solution uses k subsets, the greedy algorithm uses at most $k \ln n$ subsets.

- 1 Greedy algorithm is not too bad 😊— approximation ratio is $\ln n$
- 2 Can we find optimal solution efficiently?

Main Results

Theorem

If the optimal solution uses k subsets, the greedy algorithm uses at most $k \ln n$ subsets.

- 1 Greedy algorithm is not too bad 😊— approximation ratio is $\ln n$
- 2 Can we find optimal solution efficiently?
- 3 PCP theorem: No poly time algorithm outputs $\leq 0.999 \cdot k \ln n$ subsets unless $P = NP$

Interactive Proof Systems

- Most "famous" work in complexity theory over past decade.
- Prover claims a statement and verifier interrogates prover using randomly generated questions.



- 4 Believed that $2^{n^{1-o(1)}}$ -time is necessary to find the optimal

Proof

- ① Let n_t be # elements still not covered after t iterations of the greedy
- ② The optimal solution uses k -subsets $\Rightarrow \exists$ a subset covers n_t/k remaining elements

Proof

- ① Let n_t be # elements still not covered after t iterations of the greedy
- ② The optimal solution uses k -subsets $\Rightarrow \exists$ a subset covers n_t/k remaining elements
- ③ $n_{t+1} \leq n_t - n_t/k = n_t(1 - 1/k)$. So

$$n_t \leq n_0(1 - 1/k)^t < n_0 e^{-t/k} = n \cdot e^{-t/k},$$

where RHS is equal to 1 for $t = k \ln n$.

Proof

- 1 Let n_t be # elements still not covered after t iterations of the greedy
- 2 The optimal solution uses k -subsets $\Rightarrow \exists$ a subset covers n_t/k remaining elements
- 3 $n_{t+1} \leq n_t - n_t/k = n_t(1 - 1/k)$. So

$$n_t \leq n_0(1 - 1/k)^t < n_0 e^{-t/k} = n \cdot e^{-t/k},$$

where RHS is equal to 1 for $t = k \ln n$.

Discussion

For many computational hard problems, greedy algorithms provide the best solution in poly time.

Outline

- 1 Introduction
- 2 Revisit Interval Scheduling
- 3 Max Interval Sum
- 4 Huffman Codes
- 5 Set Cover
- 6 Scheduling

Introduction

Problem Description

Given n jobs with processing load t_1, \dots, t_n , schedule them into m identical machines to minimize the *makespan* (defined as the max load over all machines)

Introduction

Problem Description

Given n jobs with processing load t_1, \dots, t_n , schedule them into m identical machines to minimize the *makespan* (defined as the max load over all machines)

- 1 Let OPT be the optimal answer
- 2 It is NP-hard to find OPT even for 2 machines

Introduction

Problem Description

Given n jobs with processing load t_1, \dots, t_n , schedule them into m identical machines to minimize the *makespan* (defined as the max load over all machines)

- 1 Let OPT be the optimal answer
- 2 It is NP-hard to find OPT even for 2 machines
- 3 Our plan: Consider a simple greedy algorithm and prove its output $\leq 1.5 \cdot \text{OPT}$
- 4 How to design a greedy algorithm?

A Greedy Algorithm

procedure MINMAKESPAN(t_1, \dots, t_n, m)

Resort all jobs s.t. $t_1 \geq \dots \geq t_m$

Initialize $A_j = \emptyset$ as the load of machine $j \in [m]$

for $i = 1, \dots, n$ **do**

Find the least load A_j

$A_j = A_j \cup \{t_i\}$

A Greedy Algorithm

procedure MINMAKESPAN(t_1, \dots, t_n, m)

Resort all jobs s.t. $t_1 \geq \dots \geq t_m$

Initialize $A_j = \emptyset$ as the load of machine $j \in [m]$

for $i = 1, \dots, n$ **do**

Find the least load A_j

$A_j = A_j \cup \{t_i\}$

Correctness

- 1 Running Time: $O(n \log n + n \log m)$ by heap
- 2 Correctness: max-load $\leq 1.5 \cdot \text{OPT}$

Proof of Approximation

- 1 OBS 1: $\text{OPT} \geq t_1$
- 2 OBS 2: $\text{OPT} \geq \sum_{i=1}^n t_i / m$

Proof of Approximation

- 1 OBS 1: $\text{OPT} \geq t_1$
- 2 OBS 2: $\text{OPT} \geq \sum_{i=1}^n t_i / m$
- 3 Fact: Solution $\leq \sum_{i=1}^n t_i / m + t_1$ is a 2-approx. How to refine it?

Proof of Approximation

- ① OBS 1: $\text{OPT} \geq t_1$
- ② OBS 2: $\text{OPT} \geq \sum_{i=1}^n t_i / m$
- ③ Fact: Solution $\leq \sum_{i=1}^n t_i / m + t_1$ is a 2-approx. How to refine it?
- ④ Suppose the last job t_ℓ of machine A_1 has $\ell > m$ (otherwise A_1 just has 1 job)
- ⑤ Load of A_1 is $\leq \sum_{i=1}^n t_i / m + t_\ell$

Proof of Approximation

- ① OBS 1: $\text{OPT} \geq t_1$
- ② OBS 2: $\text{OPT} \geq \sum_{i=1}^n t_i/m$
- ③ Fact: Solution $\leq \sum_{i=1}^n t_i/m + t_1$ is a 2-approx. How to refine it?
- ④ Suppose the last job t_ℓ of machine A_1 has $\ell > m$ (otherwise A_1 just has 1 job)
- ⑤ Load of A_1 is $\leq \sum_{i=1}^n t_i/m + t_\ell$
- ⑥ But $\text{OPT} \geq 2t_\ell$ since $\ell > m$
- ⑦ One could further refine it to 4/3 by assuming t_n is the last *finished* job

Questions?