

Introduction to Algorithms

Lecture 9 BFS and Shortest Paths

Xue Chen

xuechen1989@ustc.edu.cn


2025 spring in



Outline


Shortest Path Prob

Given G and a pair of vertices s and t , find the shortest path $s \rightsquigarrow t$.




Shortest Path Prob

Given G and a pair of vertices s and t , find the shortest path $s \rightsquigarrow t$.




- 1 One of the most fundamental prob in CS
- 2 Many applications: Routing traffic, traveling/booking flights, experiments, ...
- 3 Various algorithms for different graphs

Overview



Consider the shortest path from a fixed source s to all other nodes


Overview



Consider the shortest path from a fixed source s to all other nodes

- 1 Unlike connected components, no difference between directed graphs and undirected graphs
- 2 Simplest case: BFS for unit-weight edges

Overview




Consider the shortest path from a fixed source s to all other nodes

- 1 Unlike connected components, no difference between directed graphs and undirected graphs
- 2 Simplest case: BFS for unit-weight edges
- 3 Dijkstra's algorithm for non-negative weights
- 4 Bellman-Ford algorithm for negative weights

Outline



Introduction

For a **unweighted** graph G and source $s \in V$, find the shortest distance from s to any other nodes



Introduction

For a **unweighted** graph G and source $s \in V$, find the shortest distance from s to any other nodes



Basic idea

Find all vertices in the 1st layer, 2nd layer, and so on

— focus on the **breadth of each layer** instead of the depth

Implementation

Instead of maintaining vertices on layers, use a **first-in first-out queue** to store all vertices in order

procedure BFS(s)

Initialize $v.d = +\infty$, $v.\pi = NIL$, $v.color = White$ for all v except s

ENQUEUE(Q, s)

while Q is not empty **do**

$u = \text{DEQUEUE}(Q)$

for each outgoing-edge (u, w) **do**

if $w.color = White$ **then**

$w.color = Black$, $w.\pi = u$, $w.d = u.d + 1$

 ENQUEUE(Q, w)

Implementation

Instead of maintaining vertices on layers, use a **first-in first-out queue** to store all vertices in order

procedure BFS(s)

Initialize $v.d = +\infty$, $v.\pi = \text{NIL}$, $v.\text{color} = \text{White}$ for all v except s

ENQUEUE(Q, s)

while Q is not empty **do**

$u = \text{DEQUEUE}(Q)$

for each outgoing-edge (u, w) **do**

if $w.\text{color} = \text{White}$ **then**

$w.\text{color} = \text{Black}$, $w.\pi = u$, $w.d = u.d + 1$

 ENQUEUE(Q, w)

Lemma 22.3 in CLRS

At any moment, the head v_h and tail v_t in Q satisfy $v_t.d \leq v_h.d + 1$.

Also, for two adjacent vertices v_i and v_{i+1} in Q , $v_i.d \leq v_{i+1}.d$.

Analysis

Running Time

$O(n + m)$ because each vertex will be added to Q at most once and each edge will be checked for at most once

Correctness: For any u , $u.d$ equals the length of the shortest path from s to u .

Analysis

Running Time

$O(n + m)$ because each vertex will be added to Q at most once and each edge will be checked for at most once

Correctness: For any u , $u.d$ equals the length of the shortest path from s to u .

Induction: If there is a length- d $s \rightsquigarrow v$ path, then $\exists u$ adjacent to v with a length- $(d - 1)$ $s \rightsquigarrow u$ path.


Discussion

- 1 Simplest algorithm of shortest paths
- 2 Crucially relies on unit-weight edges
- 3 Another algorithm to explore the graph — many differences compared to DFS
 - DFS uses a stack while BFS uses a queue
 - DFS goes depth first — very long path; but BFS is breadth first — expand one layer to the next
 - Different types of edges in DFS and BFS trees


Outline

Intro

How about weighted graphs?




⇒




Intro

How about weighted graphs?



⇒




Observation

BFS does not work any more:

Shortest path $S \rightsquigarrow 3$ is $S \rightarrow 1 \rightarrow 3$ not $S \rightarrow 3$


Basic Properties

Let $\delta(s, v)$ be the shortest distance from s to v and $w(u, v)$ denote the length of edge (u, v)



Basic Properties


Let $\delta(s, v)$ be the shortest distance from s to v and $w(u, v)$ denote the length of edge (u, v)



- ① $\delta(s, u) = \min_{\text{paths } i_0=s, \dots, i_k=u} \sum_{j=1}^k w(i_{j-1}, i_j)$
- ② $\delta(s, u) = \min_v \{ \delta(s, v) + w(v, u) \}$

Basic Properties

Let $\delta(s, v)$ be the shortest distance from s to v and $w(u, v)$ denote the length of edge (u, v)







- 1 $\delta(s, u) = \min_{\text{paths } i_0=s, \dots, i_k=u} \sum_{j=1}^k w(i_{j-1}, i_j)$
- 2 $\delta(s, u) = \min_v \{ \delta(s, v) + w(v, u) \}$
- 3 $\forall v, \delta(s, u) \leq \delta(s, v) + w(v, u)$
- 4 Lemma 24.17 in CLRS: All shortest paths from s form a tree

Dijkstra's Algorithm

Idea





Starting from source, find **shortest paths** one by one according to the length — or expand the shortest-path tree by greedy



Dijkstra's Algorithm

Idea

Starting from source, find **shortest paths** one by one according to the length — or expand the shortest-path tree by greedy



Record those vertices whose shortest path has been determined and find the next one **by greedy**

Formal Description

procedure DIJKSTRA(s)

Initialize $d[v] = +\infty$ for all v except s

$S = \emptyset$

Q is a min-heap maintaining $V - S$ according to d

while Q is not empty **do**

$u = \text{EXTRACT-MIN}(Q)$

$S = S \cup \{u\}$

for each outgoing-edge (u, v) **do**

if $d[v] > d[u] + w(u, v)$ **then**

$d[v] = d[u] + w(u, v)$ and adjust v 's position in the heap

Analysis

Running Time

The running time is $O(n \log n + m \log n)$ via a binary heap.

Analysis

Running Time

The running time is $O(n \log n + m \log n)$ via a binary heap.

- 1 Each node u will be extracted at most once
- 2 Each edge (u, v) will be checked once

Analysis

Running Time

The running time is $O(n \log n + m \log n)$ via a binary heap.

- 1 Each node u will be extracted at most once
- 2 Each edge (u, v) will be checked once
- 3 Each extraction and update of $d(v)$ will take $O(\log n)$ -time in the heap

Correctness

Lemma 24.6 in CLRS

Dijkstra's algorithm finds the shortest distance from s to any other nodes for graphs with non-negative weights.

Induction: At any moment, any $v \in S$ always has $d[v] = \delta(s, v)$

Correctness

Lemma 24.6 in CLRS

Dijkstra's algorithm finds the shortest distance from s to any other nodes for graphs with non-negative weights.

Induction: At any moment, any $v \in S$ always has $d[v] = \delta(s, v)$

- 1 $S = \emptyset$ in initialization
- 2 For contradiction, consider the **1st node u in S** with $d[u] \neq \delta(s, u)$ to S
- 3 Consider the correct shortest path from s to u

Correctness

Lemma 24.6 in CLRS

Dijkstra's algorithm finds the shortest distance from s to any other nodes for graphs with non-negative weights.

Induction: At any moment, any $v \in S$ always has $d[v] = \delta(s, v)$

- 1 $S = \emptyset$ in initialization
- 2 For contradiction, consider the **1st node u in S** with $d[u] \neq \delta(s, u)$ to S
- 3 Consider the correct shortest path from s to u
- 4 Say the 1st vertex not in S is w (could be u)
 $\Rightarrow w$'s parent node is in S
- 5 Then $d[w] = \delta(s, w)$ — we shall add w to S , which provides the contradiction


Discussion

- 1 Dijkstra's algorithm is a greedy algorithm that extends BFS from unit weights to non-negative weights
- 2 Heaps improve the running time to $O(n \log n + m \log n)$
- 3 Shortest paths form a tree
- 4 Next: How about graphs with negative weights?

Outline



Intro

How to find the shortest distance from s for a graph with negative weights?



Intro

How to find the shortest distance from s for a graph with negative weights?





Two Cases

- 1 \exists a cycle C whose sum is negative \Rightarrow no shortest path!
- 2 \nexists such a cycle

Intro

How to find the shortest distance from s for a graph with negative weights?




Two Cases

- 1 \exists a cycle C whose sum is negative \Rightarrow no shortest path!
- 2 \nexists such a cycle


In the 1st case, the shortest path is **undefined** for some nodes (or say does not exist)

Basic Idea




- 1 If \exists negative cycles, algorithm shall output \perp
- 2 Otherwise it finds all shortest paths from s

Basic Idea



- 1 If \exists negative cycles, algorithm shall output \perp
- 2 Otherwise it finds all shortest paths from s
- 3 Dijkstra's algorithm does not work b.c. of negative edges
- 4 In the above example, we need to re-put vertex 2 in Q — repeat this for many times

Basic Idea



- 1 If \exists negative cycles, algorithm shall output \perp
- 2 Otherwise it finds all shortest paths from s
- 3 Dijkstra's algorithm does not work b.c. of negative edges
- 4 In the above example, we need to re-put vertex 2 in Q — repeat this for many times
- 5 How many repetitions for each node? At most $n - 1$ times — the longest path will have $n - 1$ edges o.w. \exists negative cycles

Description of Bellman-Ford

procedure BELLMAN-FORD(s)

Initialize $d[v] = +\infty$ for all v except s

for $i = 1, \dots, n - 1$ **do**

for each edge (u, v) **do**

if $d[v] > d[u] + w(u, v)$ **then**

$d[v] = d[u] + w(u, v)$

for each edge (u, v) **do**

if $d[v] > d[u] + w(u, v)$ **then**

 Report \exists negative cycles

Otherwise claim $d[v]$ stores the shortest distance to s

Description of Bellman-Ford

procedure BELLMAN-FORD(s)

Initialize $d[v] = +\infty$ for all v except s

for $i = 1, \dots, n - 1$ **do**

for each edge (u, v) **do**

if $d[v] > d[u] + w(u, v)$ **then**

$d[v] = d[u] + w(u, v)$

for each edge (u, v) **do**

if $d[v] > d[u] + w(u, v)$ **then**

 Report \exists negative cycles

 Otherwise claim $d[v]$ stores the shortest distance to s

Running time: $O(n \cdot m)$.

Analysis

Correctness

- 1 If \exists a negative cycle, it always outputs negative cycles
- 2 Otherwise show $d[v] = \delta(s, v)$ for all v

Analysis

Correctness

- 1 If \exists a negative cycle, it always outputs negative cycles
- 2 Otherwise show $d[v] = \delta(s, v)$ for all v

Let the **length** of a path p be the number of edges on it:

$$\delta(s, v) = \min_{\text{any path } p} \text{weight}(p) = \min_{\text{length } i=1,2,\dots} \left\{ \min_{p \text{ of length } i} \text{weight}(p) \right\}.$$

Analysis

Correctness

- 1 If \exists a negative cycle, it always outputs negative cycles
- 2 Otherwise show $d[v] = \delta(s, v)$ for all v

Let the **length** of a path p be the number of edges on it:

$$\delta(s, v) = \min_{\text{any path } p} \text{weight}(p) = \min_{\text{length } i=1,2,\dots} \left\{ \min_{p \text{ of length } i} \text{weight}(p) \right\}.$$

If there is no negative cycle, we can stop at $i = n - 1$:

$$\delta(s, v) = \min_{\text{length } i \leq n-1} \left\{ \min_{p \text{ of length } i} \text{weight}(p) \right\}$$

Analysis

Correctness

- 1 If \exists a negative cycle, it always outputs negative cycles
- 2 Otherwise show $d[v] = \delta(s, v)$ for all v

Let the **length** of a path p be the number of edges on it:

$$\delta(s, v) = \min_{\text{any path } p} \text{weight}(p) = \min_{\text{length } i=1,2,\dots} \left\{ \min_{p \text{ of length } i} \text{weight}(p) \right\}.$$

If there is no negative cycle, we can stop at $i = n - 1$:

$$\delta(s, v) = \min_{\text{length } i \leq n-1} \left\{ \min_{p \text{ of length } i} \text{weight}(p) \right\}$$

After iteration i in BELLMAN-FORD, $d[v] = \min_{p \text{ of length } \leq i} \text{weight}(p)$

Summary

Various algorithm for various graphs

- 1 No difference between directed and undirected graph (except the negative edge)
- 2 BFS, Dijkstra, Bellman-Ford, . . .
- 3 Next: All pairs shortest paths!

Questions?